

K=5: Per-Dataset Routing Accuracy (Supervised Pretraining vs Joint Training)

